

```
In [111... import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
from sklearn.model_selection import train_test_split, cross_val_score
from sklearn.preprocessing import StandardScaler
from sklearn.linear_model import LogisticRegression
from sklearn.svm import SVC
from sklearn.ensemble import RandomForestClassifier
from sklearn.metrics import confusion_matrix, accuracy_score
import tensorflow.keras
from keras.models import Sequential
from keras.layers import Dense, LeakyReLU
%matplotlib inline
```

```
In [112... df = pd.read_csv('facebook_ads.csv', encoding='ISO-8859-1')
df
```

Out[112...

	Names	emails	Country	Time Spent on Site	Salary	Clicked
0	Martina Avila	cubilia.Curae.Phasellus@quisaccumsanconvallis.edu	Bulgaria	25.649648	55330.06006	0
1	Harlan Barnes	eu.dolor@diam.co.uk	Belize	32.456107	79049.07674	1
2	Naomi Rodriquez	vulputate.mauris.sagittis@ametconsectetueradip...	Algeria	20.945978	41098.60826	0
3	Jade Cunningham	malesuada@dignissim.com	Cook Islands	54.039325	37143.35536	1
4	Cedric Leach	felis.ullamcorper.viverra@egetmollislectus.net	Brazil	34.249729	37355.11276	0
...	...	...	...	...	...	...
494	Rigel	egestas.blandit.Nam@semvitaealiquam.com	Sao Tome and Principe	19.222746	44969.13495	0
495	Walter	ligula@Cumsociis.ca	Nepal	22.665662	41686.20425	0
496	Vanna	Cum.sociis.natoque@Sedmolestie.edu	Zimbabwe	35.320239	23989.80864	0
497	Pearl	penatibus.et@massanonante.com	Philippines	26.539170	31708.57054	0
498	Nell	Quisque.varius@arcuVivamussit.net	Botswana	32.386148	74331.35442	1

499 rows × 6 columns

Exeploration Datasets

In [113...

```
df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 499 entries, 0 to 498
Data columns (total 6 columns):
#   Column                Non-Null Count  Dtype
---  -
0   Names                  499 non-null   object
1   emails                 499 non-null   object
2   Country                499 non-null   object
3   Time Spent on Site     499 non-null   float64
4   Salary                 499 non-null   float64
5   Clicked                499 non-null   int64
dtypes: float64(2), int64(1), object(3)
memory usage: 23.5+ KB
```

```
In [114... df.duplicated().sum()
```

Out[114... 0

```
In [115... df.isna().sum()
```

Out[115... Names 0  
emails 0  
Country 0  
Time Spent on Site 0  
Salary 0  
Clicked 0  
dtype: int64

```
In [116... df.describe().round(2).T
```

Out[116...

	count	mean	std	min	25%	50%	75%	max
Time Spent on Site	499.0	32.92	9.10	5.0	26.43	33.20	39.11	60.0
Salary	499.0	52896.99	18989.18	20.0	38888.12	52840.91	65837.29	100000.0
Clicked	499.0	0.50	0.50	0.0	0.00	1.00	1.00	1.0

# Clean Datasets

```
In [117... def clean_outliers(df, cols):  
  
    for col in cols:  
        # Calculate IQR  
        q1 = df[col].quantile(0.25)  
        q3 = df[col].quantile(0.75)  
        iqr = q3 - q1  
  
        # Define outlier bounds  
        lower_bound = q1 - 1.5 * iqr  
        upper_bound = q3 + 1.5 * iqr  
  
        # Remove outliers  
        df = df[(df[col] >= lower_bound) & (df[col] <= upper_bound)]  
  
    return df
```

```
In [118... df= clean_outliers(df,df[['Time Spent on Site', 'Salary']])
```

```
In [119... df.drop(columns=['Names','emails', 'Country'], inplace=True)
```

```
In [120... df
```

Out[120...

	Time Spent on Site	Salary	Clicked
<b>0</b>	25.649648	55330.06006	0
<b>1</b>	32.456107	79049.07674	1
<b>2</b>	20.945978	41098.60826	0
<b>3</b>	54.039325	37143.35536	1
<b>4</b>	34.249729	37355.11276	0
...	...	...	...
<b>494</b>	19.222746	44969.13495	0
<b>495</b>	22.665662	41686.20425	0
<b>496</b>	35.320239	23989.80864	0
<b>497</b>	26.539170	31708.57054	0
<b>498</b>	32.386148	74331.35442	1

497 rows × 3 columns

## Machine Learning:

In [121... `scaler = StandardScaler()`In [122... `X = df[['Time Spent on Site', 'Salary']]`  
`y = df['Clicked']`In [123... `X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42)`In [124... `X_train = scaler.fit_transform(X_train)`  
`X_test = scaler.transform(X_test)`

## 1- Logistic Regression:

```
In [125... lr = LogisticRegression()  
scores_logistic_regression = cross_val_score(lr, X, y, cv=5, scoring='accuracy')  
print(f'Accuracy Score with Logistic Regression = {scores_logistic_regression.mean().round(2)}')
```

Accuracy Score with Logistic Regression = 0.91

## 2- Support Vector Classifier:

```
In [126... svc = SVC()  
scores_svc = cross_val_score(svc, X, y, cv=5, scoring='accuracy')  
print(f'Accuracy Score with Support Vector Classifier = {scores_svc.mean().round(2)}')
```

Accuracy Score with Support Vector Classifier = 0.81

## 3- Random Forest Classifier:

```
In [127... rfc = RandomForestClassifier()  
scores_Random_Forest_Classifier = cross_val_score(rfc, X, y, cv=5, scoring='accuracy')  
print(f'Accuracy Score with scores Random Forest Classifier = {scores_Random_Forest_Classifier.mean().round(2)}')
```





















Accuracy Score with scores Random Forest Classifier = 0.89


## Deep Learning:


```
In [128... ann = Sequential()  
ann.add(Dense(units=50)) # 1st Hidden Layer  
ann.add(LeakyReLU())  
ann.add(Dense(units=50, activation='relu')) # 2nd Hidden Layer  
ann.add(Dense(units=1, activation='sigmoid')) # Output Layer
```


```
In [129... ann.compile(optimizer = 'adam', loss = 'binary_crossentropy', metrics = ['accuracy'])
```


```
In [130... ann.fit(X_train, y_train, batch_size = 32, epochs = 100)
```


Epoch 1/100  
13/13  2s 3ms/step - accuracy: 0.5753 - loss: 0.6388  
Epoch 2/100  
13/13  0s 3ms/step - accuracy: 0.8747 - loss: 0.5088  
Epoch 3/100  
13/13  0s 3ms/step - accuracy: 0.8980 - loss: 0.4174  
Epoch 4/100  
13/13  0s 3ms/step - accuracy: 0.8838 - loss: 0.3823  
Epoch 5/100  
13/13  0s 3ms/step - accuracy: 0.8794 - loss: 0.3354  
Epoch 6/100  
13/13  0s 3ms/step - accuracy: 0.8896 - loss: 0.3117  
Epoch 7/100  
13/13  0s 3ms/step - accuracy: 0.9064 - loss: 0.2806  
Epoch 8/100  
13/13  0s 3ms/step - accuracy: 0.9009 - loss: 0.2545  
Epoch 9/100  
13/13  0s 3ms/step - accuracy: 0.9071 - loss: 0.2616  
Epoch 10/100  
13/13  0s 3ms/step - accuracy: 0.8925 - loss: 0.2774  
Epoch 11/100  
13/13  0s 3ms/step - accuracy: 0.9079 - loss: 0.2686  
Epoch 12/100  
13/13  0s 3ms/step - accuracy: 0.9069 - loss: 0.2770  
Epoch 13/100  
13/13  0s 3ms/step - accuracy: 0.9263 - loss: 0.2440  
Epoch 14/100  
13/13  0s 3ms/step - accuracy: 0.9206 - loss: 0.2511  
Epoch 15/100  
13/13  0s 3ms/step - accuracy: 0.8962 - loss: 0.2736  
Epoch 16/100  
13/13  0s 3ms/step - accuracy: 0.9219 - loss: 0.2203  
Epoch 17/100  
13/13  0s 3ms/step - accuracy: 0.9286 - loss: 0.2391  
Epoch 18/100  
13/13  0s 2ms/step - accuracy: 0.9314 - loss: 0.2359  
Epoch 19/100  
13/13  0s 2ms/step - accuracy: 0.9054 - loss: 0.2440  
Epoch 20/100  
13/13  0s 2ms/step - accuracy: 0.9218 - loss: 0.2437  
Epoch 21/100


13/13  0s 2ms/step - accuracy: 0.8696 - loss: 0.3415  
Epoch 22/100


13/13  0s 3ms/step - accuracy: 0.9162 - loss: 0.2485  
Epoch 23/100


13/13  0s 4ms/step - accuracy: 0.9266 - loss: 0.2176  
Epoch 24/100


13/13  0s 4ms/step - accuracy: 0.9141 - loss: 0.2556  
Epoch 25/100


13/13  0s 3ms/step - accuracy: 0.9057 - loss: 0.2855  
Epoch 26/100


13/13  0s 2ms/step - accuracy: 0.9036 - loss: 0.2676  
Epoch 27/100


13/13  0s 3ms/step - accuracy: 0.9236 - loss: 0.2483  
Epoch 28/100


13/13  0s 3ms/step - accuracy: 0.9283 - loss: 0.2307  
Epoch 29/100


13/13  0s 3ms/step - accuracy: 0.9055 - loss: 0.2609  
Epoch 30/100


13/13  0s 3ms/step - accuracy: 0.9195 - loss: 0.2588  
Epoch 31/100


13/13  0s 3ms/step - accuracy: 0.9122 - loss: 0.2502  
Epoch 32/100


13/13  0s 3ms/step - accuracy: 0.9103 - loss: 0.2454  
Epoch 33/100


13/13  0s 3ms/step - accuracy: 0.9054 - loss: 0.2689  
Epoch 34/100


13/13  0s 2ms/step - accuracy: 0.9098 - loss: 0.2568  
Epoch 35/100


13/13  0s 3ms/step - accuracy: 0.9079 - loss: 0.2467  
Epoch 36/100


13/13  0s 3ms/step - accuracy: 0.8975 - loss: 0.2738  
Epoch 37/100

13/13  0s 4ms/step - accuracy: 0.9165 - loss: 0.2187  
Epoch 38/100





















13/13  0s 3ms/step - accuracy: 0.9004 - loss: 0.2558  
Epoch 39/100


13/13  0s 3ms/step - accuracy: 0.8950 - loss: 0.2814  
Epoch 40/100


13/13  0s 3ms/step - accuracy: 0.8860 - loss: 0.2760  
Epoch 41/100


13/13  0s 3ms/step - accuracy: 0.8930 - loss: 0.2737





Epoch 42/100  
13/13  0s 3ms/step - accuracy: 0.8995 - loss: 0.2693  
Epoch 43/100  
13/13  0s 2ms/step - accuracy: 0.9063 - loss: 0.2561  
Epoch 44/100  
13/13  0s 3ms/step - accuracy: 0.9008 - loss: 0.2630  
Epoch 45/100  
13/13  0s 3ms/step - accuracy: 0.9200 - loss: 0.2385  
Epoch 46/100  
13/13  0s 3ms/step - accuracy: 0.9022 - loss: 0.2409  
Epoch 47/100  
13/13  0s 3ms/step - accuracy: 0.8923 - loss: 0.2755  
Epoch 48/100  
13/13  0s 3ms/step - accuracy: 0.9284 - loss: 0.2229  
Epoch 49/100  
13/13  0s 3ms/step - accuracy: 0.8780 - loss: 0.2965  
Epoch 50/100  
13/13  0s 2ms/step - accuracy: 0.9272 - loss: 0.2282  
Epoch 51/100  
13/13  0s 4ms/step - accuracy: 0.9065 - loss: 0.2433  
Epoch 52/100  
13/13  0s 2ms/step - accuracy: 0.9273 - loss: 0.2290  
Epoch 53/100  
13/13  0s 3ms/step - accuracy: 0.9317 - loss: 0.2187  
Epoch 54/100  
13/13  0s 2ms/step - accuracy: 0.9059 - loss: 0.2373  
Epoch 55/100  
13/13  0s 2ms/step - accuracy: 0.8974 - loss: 0.2712  
Epoch 56/100  
13/13  0s 3ms/step - accuracy: 0.8942 - loss: 0.2613  
Epoch 57/100  
13/13  0s 4ms/step - accuracy: 0.9081 - loss: 0.2502  
Epoch 58/100  
13/13  0s 4ms/step - accuracy: 0.8907 - loss: 0.2556  
Epoch 59/100  
13/13  0s 3ms/step - accuracy: 0.9094 - loss: 0.2387  
Epoch 60/100  
13/13  0s 3ms/step - accuracy: 0.9294 - loss: 0.2085  
Epoch 61/100  
13/13  0s 3ms/step - accuracy: 0.9143 - loss: 0.2289  
Epoch 62/100


13/13  0s 7ms/step - accuracy: 0.9141 - loss: 0.2238  
Epoch 63/100


13/13  0s 2ms/step - accuracy: 0.8956 - loss: 0.2682  
Epoch 64/100


13/13  0s 2ms/step - accuracy: 0.9025 - loss: 0.2527  
Epoch 65/100


13/13  0s 3ms/step - accuracy: 0.9160 - loss: 0.2361  
Epoch 66/100


13/13  0s 4ms/step - accuracy: 0.9016 - loss: 0.2486  
Epoch 67/100


13/13  0s 3ms/step - accuracy: 0.8892 - loss: 0.2830  
Epoch 68/100


13/13  0s 3ms/step - accuracy: 0.9109 - loss: 0.2428  
Epoch 69/100


13/13  0s 4ms/step - accuracy: 0.8982 - loss: 0.2572  
Epoch 70/100


13/13  0s 4ms/step - accuracy: 0.8893 - loss: 0.2548  
Epoch 71/100


13/13  0s 4ms/step - accuracy: 0.9045 - loss: 0.2466  
Epoch 72/100


13/13  0s 3ms/step - accuracy: 0.9138 - loss: 0.2248  
Epoch 73/100


13/13  0s 4ms/step - accuracy: 0.8871 - loss: 0.2831  
Epoch 74/100


13/13  0s 4ms/step - accuracy: 0.9088 - loss: 0.2550  
Epoch 75/100


13/13  0s 4ms/step - accuracy: 0.9157 - loss: 0.2087  
Epoch 76/100


13/13  0s 3ms/step - accuracy: 0.9042 - loss: 0.2622  
Epoch 77/100


13/13  0s 3ms/step - accuracy: 0.8938 - loss: 0.2574  
Epoch 78/100

13/13  0s 3ms/step - accuracy: 0.9004 - loss: 0.2468  
Epoch 79/100

13/13  0s 3ms/step - accuracy: 0.8915 - loss: 0.2834  
Epoch 80/100

13/13  0s 4ms/step - accuracy: 0.8978 - loss: 0.2603  
Epoch 81/100

13/13  0s 4ms/step - accuracy: 0.8978 - loss: 0.2333  
Epoch 82/100

13/13  0s 3ms/step - accuracy: 0.9060 - loss: 0.2570

```

Epoch 83/100
13/13 ————— 0s 3ms/step - accuracy: 0.9144 - loss: 0.2235
Epoch 84/100
13/13 ————— 0s 3ms/step - accuracy: 0.9090 - loss: 0.2232
Epoch 85/100
13/13 ————— 0s 3ms/step - accuracy: 0.9122 - loss: 0.2272
Epoch 86/100
13/13 ————— 0s 3ms/step - accuracy: 0.8940 - loss: 0.2712
Epoch 87/100
13/13 ————— 0s 3ms/step - accuracy: 0.8780 - loss: 0.3122
Epoch 88/100
13/13 ————— 0s 3ms/step - accuracy: 0.9033 - loss: 0.2393
Epoch 89/100
13/13 ————— 0s 3ms/step - accuracy: 0.8928 - loss: 0.2576
Epoch 90/100
13/13 ————— 0s 3ms/step - accuracy: 0.9052 - loss: 0.2412
Epoch 91/100
13/13 ————— 0s 3ms/step - accuracy: 0.8752 - loss: 0.3056
Epoch 92/100
13/13 ————— 0s 2ms/step - accuracy: 0.9043 - loss: 0.2590
Epoch 93/100
13/13 ————— 0s 3ms/step - accuracy: 0.9143 - loss: 0.2419
Epoch 94/100
13/13 ————— 0s 3ms/step - accuracy: 0.8981 - loss: 0.2559
Epoch 95/100
13/13 ————— 0s 3ms/step - accuracy: 0.9111 - loss: 0.2335
Epoch 96/100
13/13 ————— 0s 2ms/step - accuracy: 0.9109 - loss: 0.2326
Epoch 97/100
13/13 ————— 0s 3ms/step - accuracy: 0.9024 - loss: 0.2505
Epoch 98/100
13/13 ————— 0s 2ms/step - accuracy: 0.8919 - loss: 0.2682
Epoch 99/100
13/13 ————— 0s 2ms/step - accuracy: 0.9111 - loss: 0.2184
Epoch 100/100
13/13 ————— 0s 3ms/step - accuracy: 0.8941 - loss: 0.2434

```

Out[130... <keras.src.callbacks.history.History at 0x220f56d9a60>

In [131... `ann.history.history.keys()`

```
Out[131... dict_keys(['accuracy', 'loss'])
```


```
In [132... y_pred = ann.predict(X_test)
y_pred = (y_pred > 0.5).astype(int) # Convert probabilities to binary class labels
score_ann= accuracy_score(y_test, y_pred)
print(f'Accuracy Score with ANN = {score_ann}')
```

4/4 ————— 0s 23ms/step

Accuracy Score with ANN = 0.95

## Accuracy score Comparison:

```
In [133... plt.figure(figsize=(10,5))
models = ['Logistic Regression', 'SVM', 'Random Forest', 'ANN']
scores = [0.91, 0.81, 0.88, 0.95]
plt.bar(models,scores)
plt.title('Accuracy Score Comparison')
plt.ylabel('Accuracy Score')
plt.show()
```



## Summary:

- Exeploration Dataset:

Conducted thorough data cleaning, handled missing values, and detected outliers using the IQR method.

- Machine Learning Models:

Built and evaluated Logistic Regression, Support Vector Machine (SVM), and Random Forest models using cross-validation.

- Deep Learning Model (ANN):

Developed an Artificial Neural Network that achieved the highest performance among all models.

- Key Insights:
  - ANN outperformed all traditional models with higher accuracy.
  - Proper data preprocessing and feature scaling significantly impacted model performance.
  - Cross-validation provided a realistic evaluation compared to a simple train/test split.

# Thank you